

Very little data tampering is needed to manufacture statistically significant results: A Monte Carlo simulation study

Ian Hussey^{a,*}

^a *University of Bern,*

Abstract

Fabricating a statistically significant result does not require fabricating a dataset. I simulate a dishonest analyst who starts from data generated under the true null hypothesis and alters, deletes, or invents one data point at a time, always the point that most helps, re-running the test after every change until $p < .05$. I formalize this as a family of iterative algorithmically-greedy tampering strategies (dropping cases, switching condition labels, re-pairing values, duplicating cases, fabricating cases, and replacing cases) across three common designs: the two-group mean difference (Student's t-test), the Pearson correlation, and the 2x2 association (chi-square test). Each strategy runs until significance is reached or further tampering becomes structurally impossible, and the estimand is the effort required: the number of tampered data points, both absolutely and as a proportion of the sample. Two findings recur. First, effort is small in absolute terms and falls steeply in relative terms as samples grow. Median tamper counts rise only slowly with sample size, from a handful of data points at $N = 20$ to a few dozen at $N = 400$, so expressed as a share of the sample the effort drops by roughly an order of magnitude. Large studies are therefore proportionally easier to corrupt, not safer. Second, the strategy that is hardest to detect in the two-group design, because it leaves every univariate distribution untouched (condition switching), is also among the cheapest; its correlation analogue (value re-pairing) is cheaper than duplication and costs about as much as unprincipled case deletion, while being invisible to any check of either variable alone. I conclude that fabrication and falsification of results typically requires a worryingly small amount of actual data tampering.

1. Introduction

In November 2019, the pivotal trial of a treatment for ANCA-associated vasculitis was unblinded and its primary analysis run. The result was not significant ($p = .1025$). According to a subsequent letter from the US Food and Drug Administration, unblinded

*Corresponding author

Email address: ian.hussey@unibe.ch (Ian Hussey)

personnel employed by the sponsor then selected nine of the trial’s 331 patients for re-adjudication of their remission status, having first confirmed that reclassifying five of them would move the study across the significance threshold. The five were reclassified, the database was re-locked, and the re-run analysis returned $p = .0132$. That analysis was the only one submitted in support of approval, which was granted in 2021. In April 2026 the FDA proposed to withdraw it, concluding that the integrity of the trial had been compromised; the sponsor contests this and a hearing is pending ([U.S. Food and Drug Administration, Center for Drug Evaluation and Research, 2026](#)). Whatever its resolution, one feature of the case is worth dwelling on. Five records out of 331, just 1.5% of the sample, separated a failed trial from a successful one.

How often does it occur that data are simply tampered with until statistical significance is reached? Meta-analysis of the survey literature puts self-admitted research misconduct at 2.9% of researchers (95% CI [2.1, 3.8]), comprising 1.9% for fabrication and 3.3% for falsification, while 15.5% report having witnessed it in others ([Xie et al., 2021](#); see also [Fanelli, 2009](#)). When the National Survey on Research Integrity used a randomized response design to give respondents provable deniability, admitted fabrication rose to 4.3% and falsification to 4.2% ([Gopalakrishna et al., 2022](#)). From the other side of the desk, 24% of consulting biostatisticians in a US national survey had been asked, at least once in the preceding five years, to remove or alter data records to better support a research hypothesis ([Wang et al., 2018](#)). These are not the rates of a vanishing phenomenon, and each is likely a lower bound of the true prevalence.

Prevalence, though, is only half the question. The other half is how ‘effective’ such data manipulations actually are. Here, the literature is largely silent. Chambers (2017) set out a “dirty dozen” of ways to guarantee significant results even under the null hypothesis, the second of which is switching participants between conditions. Brown (2020) posted a short simulation of the efficacy of this in a blog post. Both convey the same intuition that a handful of unjustified changes to the data is enough to produce statistically significant results in many cases, but neither is a systematic assessment. The field has no systematic account of how much tampering a given strategy requires, whether that requirement grows with sample size, how the strategies compare with one another, or how large an effect they leave behind.

For the ambiguous end of the continuum of merely ‘questionable’ research practices, that gap has been filled. Stefan & Schönbrodt (2023) compiled 12 p -hacking strategies from the literature, identified the factors controlling each one’s severity, and simulated the resulting increase in false-positive rates. Their scope, as they state it, is the “grey area between good practice and outright scientific misconduct”. They stop deliberately at its far edge: discussing misreported test statistics, they note that at some point one must ask whether a practice is still questionable or has become outright fraud. As such, the literature says comparably little about the impact of indefensible research practices on false-positive rates. Of course, given a willingness to engage in an arbitrary amount of data falsification or fabrication, by definition one can obtain statistically significant results. Here, I ask a narrower question: how much data tampering, of what form, is typically needed to acquire significant results.

Quantifying the indefensible practices is worth doing because of how researchers appear to decide what to do. Entradas et al. (2026) find that perceived seriousness is by far

the strongest predictor of whether a researcher engages in a questionable practice, ahead of every career and demographic factor they measure. Equally, in John et al. (2012), deciding whether to exclude data after seeing its effect on the results was admitted by 38–43% of psychologists and rated 1.61 out of 2 for defensibility, while falsifying data was admitted by 0.6–1.7% and rated 0.16. These results are consistent with the idea of perceived norms shaping research behaviors. That is, if researchers are more likely to engage in practices they understand to be seriously problematic and are less likely to engage in ones that are seen as such. However, the boundary between these plausibly varies between individuals, research communities, and over time, creating what is effectively an Overton Window of questionable versus unacceptable research practices. Many research practices that were regarded as acceptable prior to the invention of the term *p*-hacking are now considered in many research communities to be substantially less acceptable. If perceptions of how problematic a given set of research behaviors are influence the prevalence of those behaviours, then establishing how relatively bad a given set of practices actually are, how little effort they take, and how much bias they inject, likely represents important information to dissuading their practice.

This perspective is an explicitly social-epistemological one, which understands researcher behaviors as a set of evolving community norms rather than moral acts. Equally, this perspective is concerned only with the undesirable statistical consequences of certain researcher behaviors, but is agnostic to researchers’ intent when taking them. With that being said, at a human level, it is easier to imagine a researcher who deletes a small number of data points to ‘reveal’ an effect they strongly believe is there, lying just under the surface of the data, when it is not (e.g., cases of ‘qualitative exclusion strategies’: Schiller et al. (2018)) than to imagine scientists simply making up datasets wholesale [although the latter does occur; Verfaellie & McGwin (2011)], despite both strategies producing false-positive results. Regardless of whether these two scenarios can be equated in their intent, it is an open question whether they can be equated in their outcomes, as it is generally unclear how relatively robust or fragile most results are to modifications of only a small number of data points (although see Spiess et al., 2024).

Some definition matters are worth considering, insofar as they informed which tampering strategies were simulated. The US government’s Office of Research Integrity provides influential definitions of research misconduct and its forms (Steneck, 2007). According to the ORI definitions, fabrication is “making up data or results and recording or reporting them”. This is represented by the the fabrication, duplication and replacement strategies simulated here, which add data points. Falsification is “manipulating research materials, equipment, or processes, or changing or omitting data or results such that the research is not accurately represented in the research record”. This is represented by the dropping strategy (which omits observations) and the condition switching and re-pairing strategies (which change labels and pairings without adding or deleting anything). The taxonomy that carries the regulatory and moral weight thus classifies by mechanism: was the data point invented, or was it altered or omitted? Of course, the line is not necessarily clear here, given that an altered data point is equivalent to a removed data point plus a newly added data point.

In this article I do three things. First, I formalize data tampering as a family of iterative greedy algorithms (dropping, condition switching, re-pairing, duplication, fabrication, and

replacement) applied to three of the most common elementary designs: the two-group mean difference (Student’s t -test), the Pearson correlation, and the 2×2 association (chi-square test). Second, I estimate the effort each strategy requires, with the full design specified under the ADEMP framework (Morris et al., 2019; Siepe et al., 2024) and every estimate accompanied by a Monte Carlo standard error. I invert the design used in previous simulations to understand the impact of p -hacking. Where Stefan & Schönbrodt (2023) fix a budget of manipulations and measure the resulting false-positive rate, I fix the goal (statistical significance) and measure the number of tampered data points required to reach it. I take this to be the more ecologically valid framing for deliberate tampering: an analyst who has decided to manufacture a result does not set a budget in advance and stop when it is exhausted, but continues until the p -value crosses. It also yields effort directly, in units of edited data points. Because the simulations run without any tampering budget, the complete distribution of required effort is observed, and the proportion of datasets exceeding any given budget is recovered afterwards from the same runs. Third, I show that the effort required is small, that it does not grow with sample size in proportional terms, and that the strategies hardest to detect are among the cheapest.

2. Methods

The design is specified under the ADEMP framework (Aims, Data-generating processes, Estimands, Methods, Performance measures) (Morris et al., 2019), as operationalized for psychology by Siepe et al. (2024). The three research compendia containing all simulation code, per-strategy results, and extended rationale are available in the project repository (see *Data and code availability*); this section compresses them.

2.1. Aims

The study is descriptive: there is no closed-form benchmark for the minimum number of tampered points needed to manufacture significance, because that number is a property of a tampering algorithm applied to random data, not an estimator of a population parameter. For each strategy and sample size I aim to estimate (1) how many tampered data points are required to first reach $p < .05$, absolutely and as a proportion of the sample, and how this scales with sample size; (2) the proportion of datasets requiring more than a given share of the sample (post hoc budgets of 10% and 25%), and the proportion for which significance is never attained at all; and (3) how the strategies compare on these measures.

2.2. Data-generating processes

All data are generated under the true null, so every significant result is a false positive produced entirely by the tampering.

- Mean difference. Two independent groups of `n_per_condition` observations each, $Y_{ij} \sim \mathcal{N}(\mu_j, \sigma^2)$ with $\mu_{\text{control}} = \mu_{\text{intervention}} = 0$ and $\sigma = 1$.
- Correlation. A bivariate Normal sample of n observations with $\rho = 0$ and unit variances.
- 2×2 association. Two groups (exposed, not exposed) of `n_per_condition` observations with a binary outcome drawn at the same Bernoulli prevalence in both groups; prevalence is varied (.20 vs .05) to contrast a common with a rare outcome.

Per-group sample sizes are 10 and 25 to 200 in steps of 25 (total $N = 20\text{--}400$; the correlation setting uses the same total N s). Fully crossing strategy with sample size (and prevalence) gives 45 conditions in each continuous setting and 54 in the binary setting, each replicated 10,000 times.

2.3. Estimands

The estimand is `tampers_needed`: the number of tampering operations applied before the test first returns $p < .05$, reported both as a count and as a proportion of the sample. The compendia also record the effect size of each final tampered dataset, but those are not analysed here: this article is concerned with the effort required to manufacture a false positive rather than with the size of the effect that results.

2.4. Methods (tampering strategies)

Every strategy is an iterative greedy procedure: generate a null dataset, test it, and while non-significant, apply one tampering operation targeting the data point most opposed to the desired result, then re-test. Table 1 defines the operations. All strategies push toward a positive effect by convention; tests are two-sided at $\alpha = .05$ (equal-variance t -test; Pearson correlation test; chi-square test without continuity correction).

Table 1: The tampering strategies and their operationalization in each setting. Each operation counts as one tampered data point. For binary data, duplication, fabrication, and replacement collapse onto the same atomic operation (adding a favourable cell) and are treated as a single strategy.

Strategy	Mean difference	Correlation	2×2 association
Dropping	Delete the most hostile case (lowest in intervention or highest in control)	Delete the most discordant case (most negative centred cross-product)	Delete a 1 from not-exposed or a 0 from exposed
Label switching	Move the most hostile case's <i>condition label</i> to the other group, alternating direction to keep groups balanced	<i>Re-pairing</i> : swap the y -values of two cases so the most discordant case receives the y matching its x -rank	Move a not-exposed 1 to exposed or an exposed 0 to not-exposed, alternating direction
Duplication	Duplicate a case from each group's most favourable 10%	Duplicate one of the most concordant cases	(coincides with fabrication)
Fabrication	Add invented cases: control $\sim \mathcal{N}(0, 0.2)$, intervention $\sim \mathcal{N}(5, 0.2)$	Add a high-leverage concordant pair at $\pm(3, 3)$ with small scatter	Add a 1 to exposed or a 0 to not-exposed
Replacement	Replace the most hostile case with an invented one (whichever of the two candidate replacements most lowers p)	Replace the most discordant case's y with a concordant high-leverage value	(coincides with fabrication)

The label-switching column deserves emphasis because of its forensic properties: condition switching leaves every univariate distribution untouched, and re-pairing preserves both marginal distributions *exactly* (the set of y -values is only permuted), so neither is detectable by any univariate check. Swapping *entire* conditions, as opposed to individual

cases, is handled analytically rather than simulated: it merely doubles the directional false-positive rate to α , adding $\frac{1}{2}\alpha$.

No tampering budget is imposed. Each strategy runs until significance is reached or a structural limit is hit, such removing all data under dropping. A loop guard is included at $10\times$ the sample size as a termination backstop (e.g., a one-step greedy cannot cross the two-sided test’s $p \approx 1$ valley when the initial dataset happens to show a strong effect in the *opposite* direction; such iterations, well under 1% of runs, are recorded as never attaining significance). Because the full unbudgeted counts are recorded, the proportion of datasets exceeding *any* tampering budget is computable after the fact; I report budgets of 10% and 25% of the sample.

2.5. Performance measures and Monte Carlo uncertainty

An iteration that never attains significance has no finite tamper count (i.e., is right-censored). All percentile-based summaries (median, 95% prediction interval) are therefore computed over all iterations with the never-significant ones ranked above every finite count; a percentile falling among them is reported as *never* rather than being silently replaced by the survivors’ percentile, which would answer a different question (“given that tampering worked, how many edits?”) and bias the effort downward. The mean cannot be computed this way at all (the censored counts are unbounded above), so means are reported only as conditional-on-success skew diagnostics. Two intervals are used. The interquartile range is plotted, because one condition (correlation duplication at the smallest sample size) has a 97.5th percentile at nearly nine times the sample and would otherwise set the scale for every panel. The 95% prediction interval is reported in Table 2. Effort as a proportion is expressed relative to the *original* sample size, so strategies that append cases, or re-tamper a case already tampered with, can exceed 1.

Every reported estimate carries a Monte Carlo standard error: binomial for proportions, closed-form for means, and nonparametric bootstrap for medians (Morris et al., 2019; Siepe et al., 2024). In tables, decimal places are set from the MCSE (two significant figures, with the estimate rounded to match) rather than fixed per table. Simulations were parallelized with reproducible L’Ecuyer–CMRG streams, so results are identical regardless of worker count.

3. Results

3.1. Effort is small, and falls as a share of the sample as studies grow

Across both continuous settings the absolute number of tampered data points needed to manufacture significance grows with sample size, but only slowly: median counts run from 2–10 points at $N = 20$ to 4–24 at $N = 400$, a twentyfold increase in sample size met by roughly a threefold increase in effort (Figure 1, Figure 2). Expressed as a share of the sample, effort therefore falls: from 10.0%–50.0% at $N = 20$ to 1.0%–6.0% at $N = 400$, a reduction of between 6.7- and 10.0-fold depending on the strategy. Manufacturing a false positive is not merely no harder in a large study than in a small one; proportionally it is considerably easier. The intuition that a large sample is robust to a few altered observations is exactly backwards. Framed as a budget: at $N = 400$, the proportion of datasets that would require more than 10% of the sample to be tampered is at most 21.4%




Figure 1: Mean-difference (t-test) setting: the number of tampered data points needed to first reach $p < .05$ (top) and the same effort as a proportion of the sample (bottom). Points are censoring-aware medians and bars are interquartile ranges. A bar ending in a hollow triangle is open-ended: more than a quarter of those datasets never attained significance, so the upper quartile has no finite value and the bar is drawn to the top of the data range. The lower panel expresses effort relative to the *original* sample size, so strategies that append cases or re-tamper the same case can exceed 1. Strategies are ordered by median proportion tampered. $K = 10,000$ iterations per condition.




Figure 2: Correlation setting: tampered data points needed to first reach $p < .05$ (top) and as a proportion of the sample (bottom). Points are censoring-aware medians and bars are interquartile ranges; a bar ending in a hollow triangle is open-ended, its upper quartile falling among datasets that never attained significance. Effort in the lower panel is relative to the *original* sample size, so strategies that append cases can exceed 1. $K = 10,000$ iterations per condition.




Figure 3: 2x2 association (chi-square) setting, by outcome prevalence: the number of tampered cells needed to first reach $p < .05$ (top) and the same effort as a proportion of the sample (bottom). Points are censoring-aware medians and bars are interquartile ranges; a bar ending in a hollow triangle is open-ended, its upper quartile falling among datasets that never attained significance. At prevalence .05 and $N = 20$, a majority of datasets cannot be pushed to significance at all by dropping, so even the median is not attained and no point is drawn for it. $K = 10,000$ iterations per condition.

for the least efficient continuous-setting strategy, and for the cheapest it is 0.0%. Only 21.4% of datasets resist the least effective strategy altogether at that sample size.

The 2×2 setting is the partial exception that illustrates the mechanism (Figure 3). With a rare outcome (prevalence .05) and small samples there is simply not enough signal-bearing data to tamper with: at $N = 20$, 60% of datasets cannot be pushed to significance by dropping at all as the strategy exhausts the table’s hostile cells first. Where tampering is structurally possible, rare outcomes are cheap to tip, because each fabricated or relabelled case moves the test statistic disproportionately.

3.2. The stealthiest strategies are among the cheapest

Strategies in every figure are ordered by their median proportion tampered, computed from the data at render time. The forensically critical result is where the label-switching family sits in that ordering. In the two-group design, condition switching is the second cheapest of the five strategies (median 5 switched labels at $N = 100$) despite leaving every univariate distribution untouched. Its correlation analogue, re-pairing (median 8 re-paired values at $N = 100$), is not the cheapest in its setting. High-leverage replacement and fabrication are cheaper, but it costs about the same as unprincipled case deletion and less than duplication, while preserving both marginal distributions exactly. Neither manipulation is detectable by any check of a variable in isolation; only the joint structure is disturbed. Efficiency and detectability are therefore not in tension for the would-be fraudster: in each setting, at least one of the cheapest strategies is also among the least visible.

Table 2: Tampering effort at a representative sample size (total $N = 100$), by setting and strategy. Values in parentheses are ± 1 Monte Carlo standard error. Never is the proportion of datasets for which significance was never attained, tampering having run to a structural limit. The two budget columns are post hoc tampering budgets: the proportion of datasets needing more than that share of the sample tampered, never-significant datasets exceeding every budget by definition. The median is computed over all iterations with never-significant ones ranked above every finite count, and is also given as a proportion of the sample (Prop.). PI is the 95% prediction interval, the 2.5th-97.5th percentile of the tamper counts across iterations; the figures show interquartile ranges instead, for the reason given in the Methods. D/f/r denotes the duplication, fabrication and replacement strategies, which coincide for binary data.

Setting	Strategy	Never	> 10%	> 25%	Median	95% PI	Prop.
Mean difference	Condition switching	0.000 (0.000)	0.020 (0.001)	0.000 (0.000)	5.0 (0.0)	[0, 10]	0.050 (0.000)
	Dropping	0.000 (0.000)	0.414 (0.005)	0.001 (0.000)	9.0 (0.0)	[0, 20]	0.090 (0.000)
	Duplication	0.000 (0.000)	0.659 (0.005)	0.081 (0.003)	14.0 (0.0)	[0, 28]	0.140 (0.000)
	Fabrication	0.000 (0.000)	0.625 (0.005)	0.001 (0.000)	12.0 (0.6)	[0, 22]	0.120 (0.006)
	Replacement	0.000 (0.000)	0.062 (0.002)	0.000 (0.000)	4.0 (0.0)	[0, 13]	0.040 (0.000)
Correlation	Dropping	0.000 (0.000)	0.345 (0.005)	0.001 (0.000)	8.0 (0.0)	[0, 19]	0.080 (0.000)
	Duplication	0.000 (0.000)	0.638 (0.005)	0.102 (0.003)	13.0 (0.2)	[0, 32]	0.130 (0.002)
	Fabrication	0.000 (0.000)	0.000 (0.000)	0.000 (0.000)	4.0 (0.0)	[0, 6]	0.040 (0.000)
	Re-pairing	0.233 (0.004)	0.354 (0.005)	0.233 (0.004)	8.0 (0.0)	[0, never]	0.080 (0.000)
	Replacement	0.000 (0.000)	0.000 (0.000)	0.000 (0.000)	3.0 (0.0)	[0, 6]	0.030 (0.000)
2x2, prev. 0.05	Condition switching	0.006 (0.001)	0.140 (0.003)	0.034 (0.002)	5.0 (0.0)	[0, 30]	0.050 (0.000)
	Dropping	0.057 (0.002)	0.735 (0.004)	0.445 (0.005)	25.0 (0.5)	[0, never]	0.250 (0.005)
	D/f/r	0.006 (0.001)	0.617 (0.005)	0.032 (0.002)	12.0 (0.5)	[0, 26]	0.120 (0.005)
2x2, prev. 0.2	Condition switching	0.000 (0.000)	0.263 (0.004)	0.000 (0.000)	7.0 (0.0)	[0, 15]	0.070 (0.000)
	Dropping	0.000 (0.000)	0.629 (0.005)	0.040 (0.002)	13.0 (0.0)	[0, 27]	0.130 (0.000)
	D/f/r	0.000 (0.000)	0.788 (0.004)	0.224 (0.004)	18.0 (0.2)	[0, 35]	0.180 (0.002)

4. Discussion

The distance from a null dataset to a statistically significant one is short and cheap. The median effort is a small percentage of the sample across strategies and settings, and for most strategies a minority of datasets require more than 10% of the sample to be tampered. Where tampering fails outright it is because it is structurally impossible rather than merely expensive. Specifically, the rare-outcome 2×2 cells, where too few hostile cases exist to move the table at all. The share of the sample that must be altered falls as samples grow, by roughly an order of magnitude across the range examined, so the intuition that a large study is protected by its size is not merely optimistic but inverted; this extends the finding for p -hacking that larger samples alone do not reduce its success rate (Stefan & Schönbrodt, 2023). The strategies that preserve univariate distributions (condition switching and re-pairing) are among the most effective. Further research is necessary to understand the relative statistical detectability of each of these strategies.

Adding invented cases, deleting inconvenient ones, and relabelling existing ones all require a comparable number of edits. The one clause of the definitions that unifies them is a consequence criterion rather than a mechanism one, such that the research is not accurately represented in the research record, and every strategy here satisfies it identically. Each produces a result that did not exist before.

The practical asymmetry, then, is not in the acts but in how they are regarded. Excluding inconvenient participants is admitted by 38–43% of researchers and rated as close to defensible; falsifying data is admitted by under 2% and rated as indefensible (John et al., 2012). When respondents are given provable deniability, admitted fabrication and falsification converge to indistinguishable rates (Gopalakrishna et al., 2022) — evidence that the gap is substantially one of reporting rather than of behaviour. And since perceived seriousness is the strongest available predictor of engagement (Entradas et al., 2026), a seriousness ranking that is mis-set on the dimension these results measure is not merely an intellectual inconsistency: it is a live cause of the behaviour. Deleting a hostile case and inventing a favourable one are, in effort and in damage to the evidence, the same act.

The ethics of reporting the results of this simulation study were given considerable thought. Although the simulations were designed to make a continuum visible rather than to provide a recipe, in principle, this article documents knowledge that could be weaponised by bad actors. At the same time, the tampering strategies simulated here have been publicly described for years (Brown, 2020; Chambers, 2017); what is added here is quantification. This quantification could help change future scientific behavior, by raising awareness around how impactful even a very small number of data tamperers can be.

The strategies employed in this simulation are deliberately simple and greedy. Real, detection-aware fraud could plausibly trade efficiency for plausibility. However, there are of course many cases of detection-unaware fraud caught because of impossible or implausible results and data, and the prevalence of undetected, relatively expert research fraud is unknown by its very nature. Conversely, several naive strategies here leave potentially detectable signatures (e.g., truncated distributions after dropping; leverage outliers after fabrication; shifted marginal prevalence in the 2×2). Future research is needed to examine such detectability via statistical means and modern forensic tools (Brown & Heathers, 2017; Heathers et al., 2018; Wilkinson et al., 2025). However, the other distribution-preserving

strategies imply that relatively statistically invisible tampering is cheap and effective. Auditable chain-of-custody for data, such as those employed in medical clinical trials that allowed the the US FDA to detect abnormalities in the trial, as discussed in the introduction, would be needed to prevent or detect such forms of tampering (U.S. Food and Drug Administration, Center for Drug Evaluation and Research, 2026). Infrastructure for auditable chain-of-custody is rare in other fields.

Data and code availability

All simulation code and results, including extended ADEMP specifications, per-condition tables, and Monte Carlo standard errors for every estimate, are available at <https://github.com/ianhussey/simulation-impact-of-data-tampering>.

References

- Brown, N. J. L. (2020). *In his excellent book “The 7 Deadly Sins of Psychology,” @chrisdc77 provides a list of 12 tips for people who want to get away with scientific fraud. Let’s skip the question of where he acquired that expertise and look at just one tip: Swap participants between conditions. /1.* <https://x.com/sTeamTraen/status/1329203865649094656>
- Brown, N. J. L., & Heathers, J. A. J. (2017). The GRIM Test: A Simple Technique Detects Numerous Anomalies in the Reporting of Results in Psychology. *Social Psychological and Personality Science*, 8(4), 363–369. <https://doi.org/10.1177/1948550616673876>
- Chambers, C. (2017). *The seven deadly sins of psychology: A manifesto for reforming the culture of scientific practice*. Princeton University Press. <https://doi.org/10.1515/9781400884940>
- Entradas, M., Feng, Y., & Sousa, I. C. e. (2026). The ‘shades of grey’ in research integrity—Researchers admit to questionable research practices that they do not perceive to be serious. *PLOS ONE*, 21(1), e0339056. <https://doi.org/10.1371/journal.pone.0339056>
- Fanelli, D. (2009). How many scientists fabricate and falsify research? A systematic review and meta-analysis of survey data. *PLoS ONE*, 4(5), e5738. <https://doi.org/10.1371/journal.pone.0005738>
- Gopalakrishna, G., Riet, G. ter, Vink, G., Stoop, I., Wicherts, J. M., & Bouter, L. M. (2022). Prevalence of questionable research practices, research misconduct and their potential explanatory factors: A survey among academic researchers in The Netherlands. *PLOS ONE*, 17(2), e0263023. <https://doi.org/10.1371/journal.pone.0263023>
- Heathers, J. A. J., Anaya, J., Zee, T. van der, & Brown, N. J. L. (2018). *Recovering data from summary statistics: Sample Parameter Reconstruction via Iterative TEchniques (SPRITE)*. PeerJ Preprints. <https://doi.org/10.7287/peerj.preprints.26968v1>
- John, L. K., Loewenstein, G., & Prelec, D. (2012). Measuring the prevalence of questionable research practices with incentives for truth telling. *Psychological Science*, 23(5), 524–532. <https://doi.org/10.1177/0956797611430953>
- Morris, T. P., White, I. R., & Crowther, M. J. (2019). Using simulation studies to evaluate statistical methods. *Statistics in Medicine*, 38(11), 2074–2102. <https://doi.org/10.1002/sim.8086>
- Schiller, D., Monfils, M.-H., Raio, C. M., Johnson, D. C., LeDoux, J. E., & Phelps, E. A. (2018). Addendum: Preventing the return of fear in humans using reconsolidation

- update mechanisms. *Nature*, 562(7727), E21–E21. <https://doi.org/10.1038/s41586-018-0405-7>
- Siepe, B. S., Bartoš, F., Morris, T. P., Boulesteix, A.-L., Heck, D. W., & Pawel, S. (2024). Simulation studies for methodological research in psychology: A standardized template for planning, preregistration, and reporting. *Psychological Methods*. <https://doi.org/10.1037/met0000695>
- Spiess, A.-N., Rödiger, S., Schaks, M., Burdukiewicz, M., & Tellinghuisen, J. (2024). *Scientific reasoning driven by influential data: Resuscitate dfstat!* bioRxiv. <https://doi.org/10.1101/2024.10.30.621016>
- Stefan, A. M., & Schönbrodt, F. D. (2023). Big little lies: A compendium and simulation of p-hacking strategies. *Royal Society Open Science*, 10(2), 220346. <https://doi.org/10.1098/rsos.220346>
- Steneck, N. H. (2007). *Introduction to the Responsible Conduct of Research*. <https://doi.org/10.1037/e638422011-001>
- U.S. Food and Drug Administration, Center for Drug Evaluation and Research. (2026). *Proposal to Withdraw Marketing Approval; Notice of Opportunity for a Hearing: TAVNEOS (avacopan) capsule, 10 mg, NDA 214487* (Letter to {Amgen}, {Inc}. Docket No. FDA-2026-N-1321). U.S. Food; Drug Administration. <https://www.fda.gov/media/192160/download>
- Verfaellie, M., & McGwin, J. (2011). The case of Diederik Stapel. *APA Psychological Science Agenda*. <https://www.apa.org/science/about/psa/2011/12/diederik-stapel>
- Wang, M. Q., Yan, A. F., & Katz, R. V. (2018). Researcher Requests for Inappropriate Analysis and Reporting: A U.S. Survey of Consulting Biostatisticians. *Annals of Internal Medicine*, 169(8), 554–558. <https://doi.org/10.7326/M18-1230>
- Wilkinson, J., Heal, C., Flemyng, E., Antoniou, G. A., Aburrow, T., Alfirevic, Z., Avenell, A., Barbour, V., Berghella, V., Bishop, D. V. M., Bordewijk, E. M., Brown, N. J. L., Christopher, J., Clarke, M., Dahly, D., Dennis, J., Dicker, P., Dumville, J., Frankish, H., ... Kirkham, J. J. (2025). *INSPECT-SR: A tool for assessing trustworthiness of randomised controlled trials*. medRxiv. <https://doi.org/10.1101/2025.09.03.25334905>
- Xie, Y., Wang, K., & Kong, Y. (2021). Prevalence of Research Misconduct and Questionable Research Practices: A Systematic Review and Meta-Analysis. *Science and Engineering Ethics*, 27(4), 41. <https://doi.org/10.1007/s11948-021-00314-9>